Colaberry Platform Reliability Audit

67 defects found while deploying Ellis, 43 of 51 confirmed on review, 10 fixed upstream

Building Ellis means reading the platform's code before trusting it. Every capability so far turned up defects in the platform it watches. This audit lists all 67, how serious each is, and what happened when A.M. checked them line by line.

As of 27 Oct 2026 **Prepared by** Quincy Nkwain Ninying, Forward Deployed Engineer

Reviewed by A.M. (governance); K.D. copied on urgent findings **Platform code** `ColaberryIntern, 13cb098 to bb1d3c3`

67	22	10	43 / 51	16
findings across 17 capabilities	live risks, with a real effect at runtime	already fixed upstream	reviewed findings that held as written or worse	awaiting review or not yet sent

HOW A FINDING IS MADE

Read the code, check what is already known, label it, then send it

1. Read the platform

Each capability starts by tracing the platform code it depends on, at a pinned commit.

2. Check what is known

The platform's tests, baselines and documents come first. A known item is cited, not pitched as new.

3. Label it

Live risk, latent or reporting gap, with who is affected and the file and line behind it.

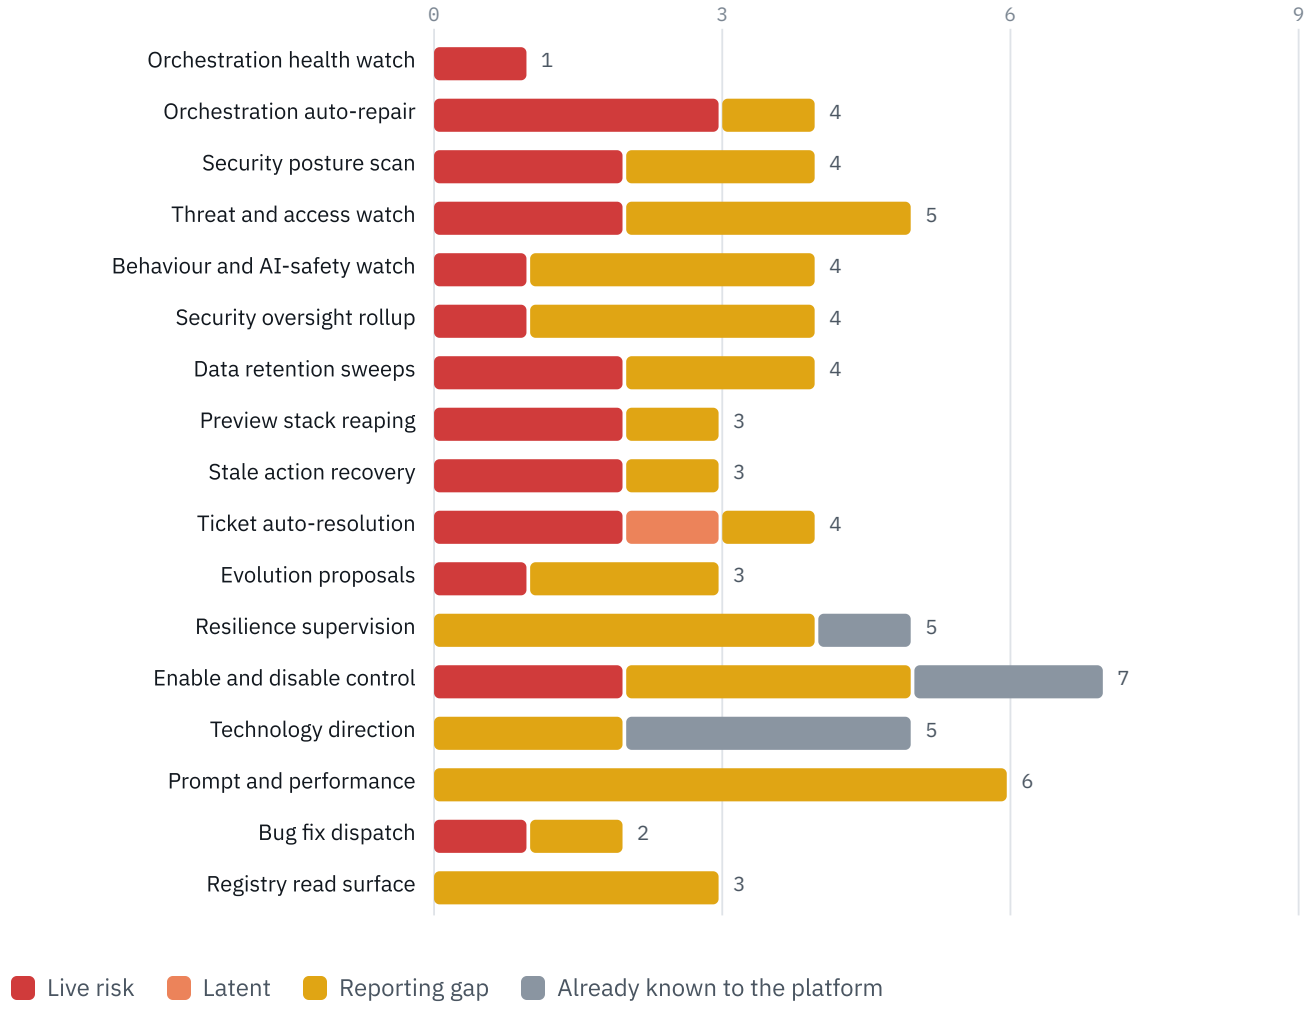
4. Send it

Security-shaped findings go to A.M. straight away, K.D. copied. Others go in labelled batches.

WHERE THE FINDINGS ARE

Findings by capability and severity

Each bar is one capability's findings, split by label. A live risk has a real effect at runtime; a reporting gap misleads a person but changes nothing; latent code cannot fire today.



HOW THEY HELD UP

Of 51 findings A.M. has reviewed, 43 held as written or worse

Five were rewritten because they did not hold as first stated, and three were already known to the platform. Each correction is in the capability write-ups.

Fixed upstream	the platform changed	10
Confirmed	held as written, or worse	27
Rewritten after review	did not hold as first stated	5
Already known	the platform had recorded it	3
Deferred upstream	waiting on a decision	6
Awaiting review	sent, no reply yet	13
Drafted, not sent	next batch	3

ALL 67 FINDINGS, BY OUTCOME



Ten defects are already fixed upstream, in two public platform changes anyone can check: eight from the first batch in pull request 3410 (branch `fix/quincy-platform-findings`), and two more that A.M. fixed on receipt in pull request 3422 (the endless retry and the duplicate send). Six are deferred until K.D. or A.M. make a decision.

LIVE RISKS

The findings with a real effect at runtime

Listed first in every message, with who is affected. Ordered by status, open items first.

#	FINDING	CAPABILITY	STATUS
49	The fix agent says a fix was applied when none was	Bug fix dispatch	Awaiting review
E3	A repair switches back on a prompt a person turned off on purpose	Orchestration auto-repair	Deferred upstream
E8	Only one security ticket can be open at a time	Security posture scan	Deferred upstream
E11	Removing a role leaves access in place for up to 12 hours	Threat and access watch	Deferred upstream
2	The grant map is documentation, not a control: writes are checked in one runtime only	Behaviour and AI-safety watch	Confirmed
8	Critical security alerts can never become tickets	Security oversight rollup	Confirmed
9	The intelligence cleanup can delete rows it never archived	Data retention sweeps	Confirmed
10	Two cleanups are single unbounded deletes holding a write lock	Data retention sweeps	Confirmed
13	A failed stop is recorded as stopped and never retried	Preview stack reaping	Confirmed
15	Reaping and archiving do not free disk; only delete does	Preview stack reaping	Confirmed
19	Switching a failing agent off closes its ticket as "recovered"	Ticket auto-resolution	Confirmed
20	A failed query or a quiet week closes lead-drop tickets as "cleared"	Ticket auto-resolution	Confirmed
25	An agent raises its own department's health score, outside its grant	Evolution proposals	Confirmed
32	Pause is ignored by five timers	Enable and disable control	Confirmed
33	A pause set during a run is undone when the run ends	Enable and disable control	Confirmed

#	FINDING	CAPABILITY	STATUS
E1	Failed agent runs are logged as "no issues" and not counted as errors	Orchestration health watch	Fixed upstream
E4	Repairs reach more rows than were flagged	Orchestration auto-repair	Fixed upstream
E5	A repair limit of 0 becomes 10	Orchestration auto-repair	Fixed upstream
E6	The secret ticket stores part of the secret	Security posture scan	Fixed upstream
E10	A staff member with a scoped role can raise themselves to owner	Threat and access watch	Fixed upstream
16	Stale recovery ignores the attempt cap, so a hung row retries forever	Stale action recovery	Fixed upstream
17	A slow row can be sent twice, to the lead who just replied	Stale action recovery	Fixed upstream

PATTERNS

Three shapes recur across the platform

Health reported that was never measured

A scanner that never ran reads normal, a failed cycle records success, a missing run reads as zero.

E1, E7, 5, 28, 43, 47

A control that is not enforced

The grant map is not checked, a guard has no caller, a pause is invisible to five timers.

2, 25, 32, 33, 50

A record that says it happened

A stop marked done when it failed, a ticket closed as recovered on no evidence, a fix announced that was never applied.

13, 19, 20, 44, 49

REGISTER

All 67 findings

E1 to E14 are the first batch, numbered before the series started; 1 to 53 follow in the order they were found.

#	FINDING	CAPABILITY	LABEL	OUTCOME
E1	Failed agent runs are logged as "no issues" and not counted as errors	Orchestration health watch	Live risk	Fixed upstream
E2	"Repairs Today" never counts the repair agent's repairs	Orchestration auto-repair	Reporting gap	Fixed upstream
E3	A repair switches back on a prompt a person turned off on purpose	Orchestration auto-repair	Live risk	Deferred upstream
E4	Repairs reach more rows than were flagged	Orchestration auto-repair	Live risk	Fixed upstream
E5	A repair limit of 0 becomes 10	Orchestration auto-repair	Live risk	Fixed upstream
E6	The secret ticket stores part of the secret	Security posture scan	Live risk	Fixed upstream
E7	Two scanners report clean after scanning nothing	Security posture scan	Reporting gap	Fixed upstream
E8	Only one security ticket can be open at a time	Security posture scan	Live risk	Deferred upstream
E9	A failed dependency audit drops out silently	Security posture scan	Reporting gap	Fixed upstream
E10	A staff member with a scoped role can raise themselves to owner	Threat and access watch	Live risk	Fixed upstream
E11	Removing a role leaves access in place for up to 12 hours	Threat and access watch	Live risk	Deferred upstream

#	FINDING	CAPABILITY	LABEL	OUTCOME
E12	The access control guardian never checks anyone's access	Threat and access watch	Reporting gap	Deferred upstream
E13	New admin logins leave no audit trail	Threat and access watch	Reporting gap	Deferred upstream
E14	The threat monitor does not look at failed logins	Threat and access watch	Reporting gap	Deferred upstream
1	No watcher compares agent actions with their grants	Behaviour and AI-safety watch	Reporting gap	Confirmed
2	The grant map is documentation, not a control: writes are checked in one runtime only	Behaviour and AI-safety watch	Live risk	Confirmed
3	"Stuck" alerts are measured from the wrong time, fleet-wide	Behaviour and AI-safety watch	Reporting gap	Confirmed
4	Error spikes on disabled agents are labelled "Unknown"	Behaviour and AI-safety watch	Reporting gap	Confirmed
5	The security director says "normal" over scanners that never ran	Security oversight rollup	Reporting gap	Confirmed
6	"Active Threats (24h)" covers about the last 80 minutes	Security oversight rollup	Reporting gap	Confirmed
7	"Scans Today" counts every security event, not scans	Security oversight rollup	Reporting gap	Confirmed
8	Critical security alerts can never become tickets	Security oversight rollup	Live risk	Confirmed
9	The intelligence cleanup can delete rows it never archived	Data retention sweeps	Live risk	Confirmed
10	Two cleanups are single unbounded deletes holding a write lock	Data retention sweeps	Live risk	Confirmed
11	The written window for generated cards does not match the code	Data retention sweeps	Reporting gap	Confirmed
12	The retention policy disagrees with the scheduler and never runs	Data retention sweeps	Reporting gap	Rewritten after review
13	A failed stop is recorded as stopped and never retried	Preview stack reaping	Live risk	Confirmed
14	Idle is measured from stack age, and the reaped set is not kept	Preview stack reaping	Reporting gap	Confirmed
15	Reaping and archiving do not free disk; only delete does	Preview stack reaping	Live risk	Confirmed
16	Stale recovery ignores the attempt cap, so a hung row retries forever	Stale action recovery	Live risk	Fixed upstream
17	A slow row can be sent twice, to the lead who just replied	Stale action recovery	Live risk	Fixed upstream
18	Recovery keeps no record, and its tests never call it	Stale action recovery	Reporting gap	Confirmed
19	Switching a failing agent off closes its ticket as "recovered"	Ticket auto-resolution	Live risk	Confirmed
20	A failed query or a quiet week closes lead-drop tickets as "cleared"	Ticket auto-resolution	Live risk	Confirmed
21	A planned reset would wipe the evidence the resolver checks	Ticket auto-resolution	Latent	Rewritten after review
22	A close and its evidence are two separate writes	Ticket auto-resolution	Reporting gap	Confirmed
23	No architect proposal can be accepted, deferred or rejected in the app	Evolution proposals	Reporting gap	Confirmed
24	No idea can ever be proposed again	Evolution proposals	Reporting gap	Rewritten after review
25	An agent raises its own department's health score, outside its grant	Evolution proposals	Live risk	Confirmed

#	FINDING	CAPABILITY	LABEL	OUTCOME
26	The resilience supervisor's header names agents that do not exist	Resilience supervision	Reporting gap	Rewritten after review
27	One hour of silence is allowed for every agent, whatever its schedule	Resilience supervision	Reporting gap	Confirmed
28	Supervisors' own rows read healthy when their reports say unknown	Resilience supervision	Reporting gap	Rewritten after review
29	Reading the circuit status can trip or silently reset a breaker	Resilience supervision	Reporting gap	Confirmed
30	Dara's three timers share one counter	Resilience supervision	Already known	Already known
31	Timer names differ from the row their switch reads	Enable and disable control	Already known	Already known
32	Pause is ignored by five timers	Enable and disable control	Live risk	Confirmed
33	A pause set during a run is undone when the run ends	Enable and disable control	Live risk	Confirmed
34	Dara's directors have two switches, and the comment says one	Enable and disable control	Reporting gap	Confirmed
35	A disable fails open during a database fault, by design	Enable and disable control	Already known	Already known
36	Most skips are not structured events	Enable and disable control	Reporting gap	Confirmed
37	A control call is audited with no record of what changed	Enable and disable control	Reporting gap	Confirmed
38	The Technology Director flags the most-erroring agent, not direction	Technology direction	Already known	Awaiting review
39	A closed task comes back within 15 minutes	Technology direction	Reporting gap	Awaiting review
40	One open task hides every other unhealthy agent	Technology direction	Reporting gap	Awaiting review
41	Most of the fleet's tooling is not recorded	Technology direction	Already known	Awaiting review
42	No written technology direction exists	Technology direction	Already known	Awaiting review
43	A loop whose every step failed is recorded as a success	Prompt and performance	Reporting gap	Awaiting review
44	Experiments change nothing and are still logged as adopted	Prompt and performance	Reporting gap	Awaiting review
45	Prompt proposals are three fixed sentences	Prompt and performance	Reporting gap	Awaiting review
46	The low-confidence check never fires	Prompt and performance	Reporting gap	Awaiting review
47	The loop measures no cost, and no runs averages to zero	Prompt and performance	Reporting gap	Awaiting review
48	The one-minute prompt monitor checks references, not behaviour	Prompt and performance	Reporting gap	Awaiting review
49	The fix agent says a fix was applied when none was	Bug fix dispatch	Live risk	Awaiting review
50	The risk-tier guard is never called outside its tests	Bug fix dispatch	Reporting gap	Awaiting review
51	The registry's next run is off by the UTC offset	Registry read surface	Reporting gap	Drafted, not sent
52	Next run is shown for agents that will not run	Registry read surface	Reporting gap	Drafted, not sent
53	A director reads "Never" while its gate runs daily	Registry read surface	Reporting gap	Drafted, not sent

SOURCES

- Capability write-ups: each finding's evidence, file and line, and any correction after review
- No evidence is not health: the rule behind the first pattern

- Platform pull request 3410: the fixes for E1, E2, E4, E5, E6, E7, E9 and E10, merged as 3ffea7c
- Platform pull request 3422: the fixes for findings 16 and 17, merged as 7c70f0a
- Colaberry platform repository: every cited line, at the commit named in each write-up