# Colaberry Platform Reliability Audit

### 82 defects found while deploying Ellis, 61 of 76 confirmed on review, 14 fixed upstream

Building Ellis means reading the platform's code before trusting it. Every capability built up to 30 October turned up defects in the platform it watches. This audit lists all 82, how serious each is, and what happened when A.M. checked them line by line.

**The register closed at 82 on 30 October.** K.D. and A.M. decided that defects are for Ellis to find, not for its builder to send for fixing. Findings 63 to 68 had been drafted and stay here as written, never sent and never reviewed. Defects found since then appear only in Ellis's output and its capability write-ups.

**As of** 2 Nov 2026 **Prepared by** Quincy Nkwain Ninying, Forward Deployed Engineer

**Reviewed by** A.M. (governance); K.D. copied on urgent findings **Platform code** `ColaberryIntern, 13cb098 to 4443087`

| 82 | 25 | 14 | 61 / 76 | 6 |
|---|---|---|---|---|
| findings across 21 capabilities | live risks, with a real effect at runtime | already fixed upstream | reviewed findings that held as written or worse | drafted and never sent: the register closed on 30 Oct |

HOW A FINDING IS MADE

## Read the code, check what is already known, label it, then send it

**1. Read the platform**

Each capability starts by tracing the platform code it depends on, at a pinned commit.

**2. Check what is known**

The platform's tests, baselines and documents come first. A known item is cited, not pitched as new.

**3. Label it**

Live risk, latent or reporting gap, with who is affected and the file and line behind it.

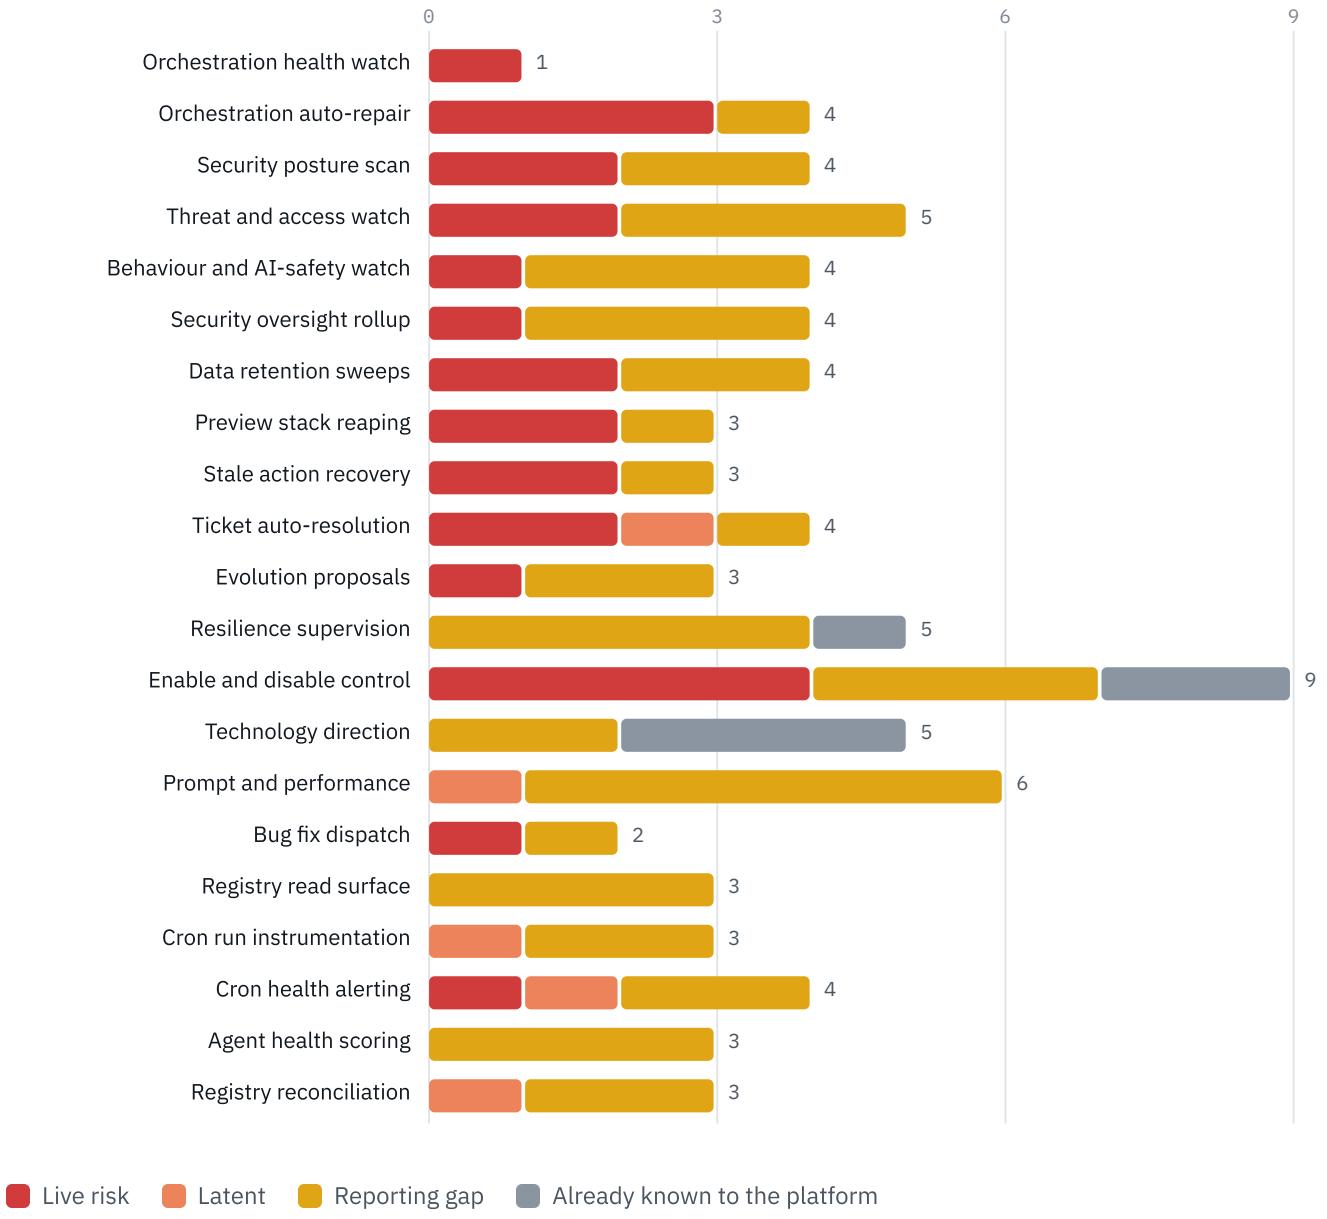
**4. Send it (until 30 Oct)**

Security-shaped findings went to A.M. straight away, K.D. copied; others in labelled batches. Since 30 Oct, Ellis reports them instead.

WHERE THE FINDINGS ARE

## Findings by capability and severity

Each bar is one capability's findings, split by label. A live risk has a real effect at runtime; a reporting gap misleads a person but changes nothing; latent code cannot fire today.



HOW THEY HELD UP

## Of 76 findings A.M. has reviewed, 61 held as written or worse

Nine were rewritten because they did not hold as first stated, and six were already known to the platform. Two were combined into one on review (39 and 40). Each correction is in the capability write-ups.

| | | |
|---|---|---|
| Fixed upstream | the platform changed | 14 |
| Confirmed | held as written, or worse | 41 |
| Rewritten after review | did not hold as first stated | 9 |
| Already known | the platform had recorded it | 6 |
| Deferred upstream | waiting on a decision | 6 |
| Awaiting review | sent, no reply yet | 0 |
| Not sent | register closed 30 Oct | 6 |

ALL 82 FINDINGS, BY OUTCOME



Fourteen defects are already fixed upstream, in five public platform changes anyone can check: eight from the first batch in pull request 3410 (branch `fix/quincy-platform-findings`), and two more that A.M. fixed on receipt in pull request 3422 (the endless retry and the duplicate send), the two pause defects in pull request 3430, one shared control check for every run wrapper, and the false "fix applied" in pull request 3435 (49). The 3430 fix missed its failure path, which still undid a pause set mid-run; that was fixed in pull request 3439 (58). Six are deferred until K.D. or A.M. make a decision.

LIVE RISKS

## The findings with a real effect at runtime

Listed first in every message, with who is affected. Ordered by status, open items first.

| # | FINDING | CAPABILITY | STATUS |
|---|---|---|---|
| E3 | A repair switches back on a prompt a person turned off on purpose | Orchestration auto-repair | Deferred upstream |
| E8 | Only one security ticket can be open at a time | Security posture scan | Deferred upstream |
| E11 | Removing a role leaves access in place for up to 12 hours | Threat and access watch | Deferred upstream |
| 2 | The grant map is documentation, not a control: writes are checked in one runtime only | Behaviour and AI-safety watch | Confirmed |
| 8 | Critical security alerts can never become tickets | Security oversight rollup | Confirmed |
| 9 | The intelligence cleanup can delete rows it never archived | Data retention sweeps | Confirmed |
| 10 | Two cleanups are single unbounded deletes holding a write lock | Data retention sweeps | Confirmed |
| 13 | A failed stop is recorded as stopped and never retried | Preview stack reaping | Confirmed |
| 15 | Reaping and archiving do not free disk; only delete does | Preview stack reaping | Confirmed |
| 19 | Switching a failing agent off closes its ticket as "recovered" | Ticket auto-resolution | Confirmed |
| 20 | A failed query or a quiet week closes lead-drop tickets as "cleared" | Ticket auto-resolution | Confirmed |
| 25 | An agent raises its own department's health score, outside its grant | Evolution proposals | Confirmed |
| 57 | WeeklyReport's Sunday timer has fired into nothing for 45 days | Enable and disable control | Confirmed |

| # | FINDING | CAPABILITY | STATUS |
|---|---|---|---|
| 61 | A schedule with uneven gaps is judged against its next gap only: a live LinkedIn job reads critical every weekend | Cron health alerting | Confirmed |
| E1 | Failed agent runs are logged as "no issues" and not counted as errors | Orchestration health watch | Fixed upstream |
| E4 | Repairs reach more rows than were flagged | Orchestration auto-repair | Fixed upstream |
| E5 | A repair limit of 0 becomes 10 | Orchestration auto-repair | Fixed upstream |
| E6 | The secret ticket stores part of the secret | Security posture scan | Fixed upstream |
| E10 | A staff member with a scoped role can raise themselves to owner | Threat and access watch | Fixed upstream |
| 16 | Stale recovery ignores the attempt cap, so a hung row retries forever | Stale action recovery | Fixed upstream |
| 17 | A slow row can be sent twice, to the lead who just replied | Stale action recovery | Fixed upstream |
| 32 | Pause is ignored by five timers | Enable and disable control | Fixed upstream |
| 33 | A pause set during a run is undone when the run ends | Enable and disable control | Fixed upstream |
| 49 | The fix agent says a fix was applied when none was, and the ticket moves to review | Bug fix dispatch | Fixed upstream |
| 58 | A pause set during a run that fails is still undone | Enable and disable control | Fixed upstream |

PATTERNS

## Three shapes recur across the platform

### Health reported that was never measured

A scanner that never ran reads normal, a failed cycle records success, a missing run reads as zero.

E1, E7, 5, 28, 43, 47

### A control that is not enforced

The grant map is not checked, a guard has no caller, a pause is undone when a run fails, a timer with no row has no control at all.

2, 25, 50, 57, 58

### Work reported that was not done

A fix announced that was never applied leads this family: it is the one a person reads and acts on. Then a ticket closed as recovered on no evidence, a stop marked done when it failed, experiments logged as adopted.

49, 19, 20, 13, 44

REGISTER

## All 82 findings

E1 to E14 are the first batch, numbered before the series started; 1 to 62 follow in the order they were found.

| # | FINDING | CAPABILITY | LABEL | OUTCOME |
|---|---|---|---|---|
| E1 | Failed agent runs are logged as "no issues" and not counted as errors | Orchestration health watch | Live risk | Fixed upstream |
| E2 | "Repairs Today" never counts the repair agent's repairs | Orchestration auto-repair | Reporting gap | Fixed upstream |
| E3 | A repair switches back on a prompt a person turned off on purpose | Orchestration auto-repair | Live risk | Deferred upstream |
| E4 | Repairs reach more rows than were flagged | Orchestration auto-repair | Live risk | Fixed upstream |

| # | FINDING | CAPABILITY | LABEL | OUTCOME |
|---|---|---|---|---|
| E5 | A repair limit of 0 becomes 10 | Orchestration auto-repair | Live risk | Fixed upstream |
| E6 | The secret ticket stores part of the secret | Security posture scan | Live risk | Fixed upstream |
| E7 | Two scanners report clean after scanning nothing | Security posture scan | Reporting gap | Fixed upstream |
| E8 | Only one security ticket can be open at a time | Security posture scan | Live risk | Deferred upstream |
| E9 | A failed dependency audit drops out silently | Security posture scan | Reporting gap | Fixed upstream |
| E10 | A staff member with a scoped role can raise themselves to owner | Threat and access watch | Live risk | Fixed upstream |
| E11 | Removing a role leaves access in place for up to 12 hours | Threat and access watch | Live risk | Deferred upstream |
| E12 | The access control guardian never checks anyone's access | Threat and access watch | Reporting gap | Deferred upstream |
| E13 | New admin logins leave no audit trail | Threat and access watch | Reporting gap | Deferred upstream |
| E14 | The threat monitor does not look at failed logins | Threat and access watch | Reporting gap | Deferred upstream |
| 1 | No watcher compares agent actions with their grants | Behaviour and AI-safety watch | Reporting gap | Confirmed |
| 2 | The grant map is documentation, not a control: writes are checked in one runtime only | Behaviour and AI-safety watch | Live risk | Confirmed |
| 3 | "Stuck" alerts are measured from the wrong time, fleet-wide | Behaviour and AI-safety watch | Reporting gap | Confirmed |
| 4 | Error spikes on disabled agents are labelled "Unknown" | Behaviour and AI-safety watch | Reporting gap | Confirmed |
| 5 | The security director says "normal" over scanners that never ran | Security oversight rollup | Reporting gap | Confirmed |
| 6 | "Active Threats (24h)" covers about the last 80 minutes | Security oversight rollup | Reporting gap | Confirmed |
| 7 | "Scans Today" counts every security event, not scans | Security oversight rollup | Reporting gap | Confirmed |
| 8 | Critical security alerts can never become tickets | Security oversight rollup | Live risk | Confirmed |
| 9 | The intelligence cleanup can delete rows it never archived | Data retention sweeps | Live risk | Confirmed |
| 10 | Two cleanups are single unbounded deletes holding a write lock | Data retention sweeps | Live risk | Confirmed |
| 11 | The written window for generated cards does not match the code | Data retention sweeps | Reporting gap | Confirmed |
| 12 | The retention policy disagrees with the scheduler and never runs | Data retention sweeps | Reporting gap | Rewritten after review |
| 13 | A failed stop is recorded as stopped and never retried | Preview stack reaping | Live risk | Confirmed |
| 14 | Idle is measured from stack age, and the reaped set is not kept | Preview stack reaping | Reporting gap | Confirmed |
| 15 | Reaping and archiving do not free disk; only delete does | Preview stack reaping | Live risk | Confirmed |
| 16 | Stale recovery ignores the attempt cap, so a hung row retries forever | Stale action recovery | Live risk | Fixed upstream |
| 17 | A slow row can be sent twice, to the lead who just replied | Stale action recovery | Live risk | Fixed upstream |
| 18 | Recovery keeps no record, and its tests never call it | Stale action recovery | Reporting gap | Confirmed |
| 19 | Switching a failing agent off closes its ticket as "recovered" | Ticket auto-resolution | Live risk | Confirmed |

| # | FINDING | CAPABILITY | LABEL | OUTCOME |
|---|---|---|---|---|
| 20 | A failed query or a quiet week closes lead-drop tickets as "cleared" | Ticket auto-resolution | Live risk | Confirmed |
| 21 | A planned reset would wipe the evidence the resolver checks | Ticket auto-resolution | Latent | Rewritten after review |
| 22 | A close and its evidence are two separate writes | Ticket auto-resolution | Reporting gap | Confirmed |
| 23 | No architect proposal can be accepted, deferred or rejected in the app | Evolution proposals | Reporting gap | Confirmed |
| 24 | No idea can ever be proposed again | Evolution proposals | Reporting gap | Rewritten after review |
| 25 | An agent raises its own department's health score, outside its grant | Evolution proposals | Live risk | Confirmed |
| 26 | The resilience supervisor's header names agents that do not exist | Resilience supervision | Reporting gap | Rewritten after review |
| 27 | One hour of silence is allowed for every agent, whatever its schedule | Resilience supervision | Reporting gap | Confirmed |
| 28 | Supervisors' own rows read healthy when their reports say unknown | Resilience supervision | Reporting gap | Rewritten after review |
| 29 | Reading the circuit status can trip or silently reset a breaker | Resilience supervision | Reporting gap | Confirmed |
| 30 | Dara's three timers share one counter | Resilience supervision | Already known | Already known |
| 31 | Timer names differ from the row their switch reads | Enable and disable control | Already known | Already known |
| 32 | Pause is ignored by five timers | Enable and disable control | Live risk | Fixed upstream |
| 33 | A pause set during a run is undone when the run ends | Enable and disable control | Live risk | Fixed upstream |
| 34 | Dara's directors have two switches, and the comment says one | Enable and disable control | Reporting gap | Confirmed |
| 35 | A disable fails open during a database fault, by design | Enable and disable control | Already known | Already known |
| 36 | Most skips are not structured events | Enable and disable control | Reporting gap | Confirmed |
| 37 | A control call is audited with no record of what changed | Enable and disable control | Reporting gap | Confirmed |
| 38 | The Technology Director flags the most-erroring agent, not direction | Technology direction | Already known | Already known |
| 39 | The Technology Director is pinned to the agent with the most lifetime errors | Technology direction | Reporting gap | Confirmed |
| 40 | Combined with 39: it is blind to every other unhealthy agent | Technology direction | Reporting gap | Confirmed |
| 41 | Most of the fleet's tooling is not recorded | Technology direction | Already known | Already known |
| 42 | No written technology direction exists | Technology direction | Already known | Already known |
| 43 | A loop whose every step failed is recorded as a success | Prompt and performance | Reporting gap | Confirmed |
| 44 | The experiment subsystem is dead and still logs adoptions | Prompt and performance | Latent | Rewritten after review |
| 45 | Prompt proposals are three fixed sentences | Prompt and performance | Reporting gap | Confirmed |
| 46 | The low-confidence check is absent, not weak (like 8) | Prompt and performance | Reporting gap | Rewritten after review |
| 47 | The loop measures no cost, and no runs averages to zero | Prompt and performance | Reporting gap | Confirmed |
| 48 | The one-minute prompt monitor checks references, not behaviour | Prompt and performance | Reporting gap | Confirmed |

| # | FINDING | CAPABILITY | LABEL | OUTCOME |
|---|---|---|---|---|
| 49 | The fix agent says a fix was applied when none was, and the ticket moves to review | Bug fix dispatch | Live risk | Fixed upstream |
| 50 | The risk-tier guard is never called outside its tests | Bug fix dispatch | Reporting gap | Confirmed |
| 51 | For the 21 Chicago-timed crons, the registry's next run is off by the UTC offset (shown on the directors' rows, not Dara's) | Registry read surface | Reporting gap | Rewritten after review |
| 52 | Next run is shown for agents that will not run | Registry read surface | Reporting gap | Confirmed |
| 53 | A director reads "Never" while its gate runs daily (stands only if Dara is switched on) | Registry read surface | Reporting gap | Confirmed |
| 54 | Four timers that fail every run read as healthy and busy: no alert, no error count | Cron run instrumentation | Reporting gap | Confirmed |
| 55 | A failed status write is recorded as a failed run that never started | Cron run instrumentation | Latent | Confirmed |
| 56 | The dead-letter queue never sees five timers fail, and nothing reads it (Dara's half stands only if she is switched on) | Cron run instrumentation | Reporting gap | Confirmed |
| 57 | WeeklyReport's Sunday timer has fired into nothing for 45 days | Enable and disable control | Live risk | Confirmed |
| 58 | A pause set during a run that fails is still undone | Enable and disable control | Live risk | Fixed upstream |
| 59 | A job reads as stopped every weekday afternoon and night: its row says every 2 hours, it runs once a weekday | Cron health alerting | Reporting gap | Rewritten after review |
| 60 | Dara's three daily duties do not run while her row is off, and three separate settings keep them out of cron alerts | Cron health alerting | Latent | Confirmed |
| 61 | A schedule with uneven gaps is judged against its next gap only: a live LinkedIn job reads critical every weekend | Cron health alerting | Live risk | Confirmed |
| 62 | A cron alert stays open after the job recovers, and an escalation opens a second one | Cron health alerting | Reporting gap | Confirmed |
| 63 | The health score means different things for different wrappers: a crash counts as an error but not a run | Agent health scoring | Reporting gap | Not sent |
| 64 | The registry page shows every agent as a green 100 when the health request fails | Agent health scoring | Reporting gap | Not sent |
| 65 | Nothing orders agents by health score, so worst-first triage means scanning every badge | Agent health scoring | Reporting gap | Not sent |
| 66 | The registry reconciler counts 24 agents twice, as phantom and as untracked, and they are neither | Registry reconciliation | Reporting gap | Not sent |
| 67 | Fourteen timers fire with no name, three of them alerting jobs, so nothing can see whether they ran | Registry reconciliation | Reporting gap | Not sent |
| 68 | The reconciler leaves out the domain and governance rows | Registry reconciliation | Latent | Not sent |

SOURCES

- Capability write-ups: each finding's evidence, file and line, and any correction after review
- No evidence is not health: the rule behind the first pattern

- Platform pull request 3410: the fixes for E1, E2, E4, E5, E6, E7, E9 and E10, merged as 3ffea7c
- Platform pull request 3422: the fixes for findings 16 and 17, merged as 7c70f0a
- Platform pull request 3430: the fixes for findings 32 and 33, merged as c699e58
- Platform pull request 3435: the fix for finding 49, merged before 4443087
- Platform pull request 3439: the fix for finding 58, merged as 7d582e8
- Colaberry platform repository: every cited line, at the commit named in each write-up