Ellis Deployment Risk Review

How an AI agent was deployed into a live production platform without the power to harm it, and what it found that could.

As of 1 Nov 2026 **Prepared by** Quincy Nkwain Ninying, Forward Deployed Engineer **Accountable** K.D. **Governance** A.M.

Scope 35 verified capabilities; platform code to `bbde6ef`

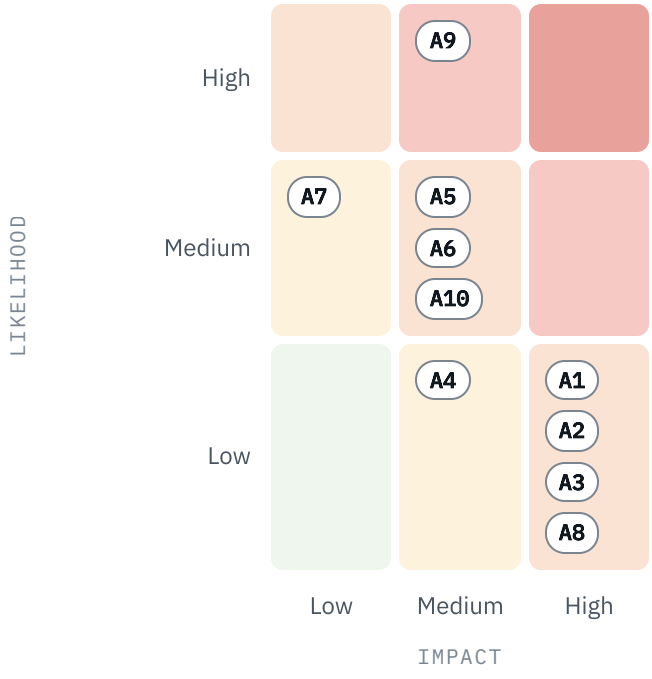
0	9	10	24	11
ways for Ellis to write to the live platform	hard limits that hold at every autonomy level	risks from the agent, each with a named control	live risks Ellis found in the platform	of those 24 already fixed upstream

PART 1 · RISK FROM THE AGENT

HEAT MAP

No risk from Ellis sits in the high-likelihood, high-impact corner

Each risk is placed after its control. The most likely one is the honest one: Ellis has not yet run on live records, so real data may surprise it.



How to read it. Rows are how likely the risk is once its control is in place; columns are how much harm it would do. Darker cells need more attention.

The high-impact column is all low likelihood. Writing to the platform, a false green, a leak and self-promotion are each blocked by construction or by a test, not by good intentions.

The medium-impact, medium-likelihood risks are about trust and judgement: a finding that overstates, a rule copied from platform code that drifts as the code moves, and a proposal followed without checking.

CONTROL REGISTER

Every risk from the agent has a control you can point to

#	RISK	CONTROL	RESIDUAL
A1	Ellis changes the live platform	Capabilities receive a read-only source with no write method; repairs run only on a sandbox copy; there is no production target in the code `src/readSource.ts, src/sandbox/sandboxStore.ts`	Low
A2	Ellis reports healthy when it cannot see (a false green)	Three answers, with unknown above clean; each capability's acceptance test removes the signal and fails if the answer is not unknown; key rules are broken on purpose to prove a test catches it `no-evidence-is-not-health.md; each capability's tests`	Low
A3	Ellis leaks a secret or personal data	Access snapshots with password, token or secret columns are refused unread; security findings pass a field whitelist; artifacts are scanned before they are published `src/access/accessSnapshot.ts, src/capabilities/securityPostureScan.ts`	Low
A4	A baseline test reaches outside its sandbox	Tests run under Node's permission model: their own folder only, no other programs, no network; a refused write is never accepted as a result; file names are checked `src/sandbox/baselineRunner.ts`	Low
A5	A finding overstates the problem and costs the owners' trust	Check the platform's tests, baselines and documents first; label each finding live, latent or reporting gap; A.M. reviewed line by line until 30 Oct, when defects moved to Ellis's own output. Of 51 reviewed, 5 were rewritten `colaberry-platform-reliability-audit.pdf`	Medium
A6	A rule Ellis copied from platform code drifts as the code changes	Every copied rule names its commit and file and line; citations are rechecked against current code when a capability touches them	Medium
A7	An export is incomplete or malformed	Malformed input is refused with exit code 1; anything missing is reported as unknown or not recorded, never filled in	Low
A8	Ellis raises its own authority	Nothing in Ellis can change an autonomy level or a policy; K.D. approves each step; 23 of 48 decisions are reserved to a person at every level `autonomyLadder.ts:213-219`	Low
A9	Real records behave differently from the tests and examples	Every answer cites its rows, so a wrong one can be traced; the first live check compares one export with Ellis's answers, with K.D.	Medium
A10	A person follows a proposal to replay or abandon without checking it	Each proposal carries its evidence and row ids; replay is labelled unproven on every answer; a person decides at every level `src/capabilities/deadLetterTriage.ts`	Medium

PART 2 · RISK ELLIS FOUND IN THE PLATFORM

LIVE RISKS BY AREA

24 live risks across five areas; 10 are already fixed

From the Colaberry Platform Reliability Audit. Each cell counts the area's live risks by status.

AREA	FIXED	CONFIRMED	DEFERRED	AWAITING	TOTAL
Security and access	**2**	**1**	**2**	0	**5**

AREA	FIXED	CONFIRMED	DEFERRED	AWAITING	TOTAL
a staff role raised to owner; a secret kept in its ticket; critical alerts never ticketed (E6, E10, E8, E11, 8)					
Control and containment the grant map is not enforced; a pause is still undone when a run fails (2, 25, 32, 33, 58, E3, E4, E5)	**5**	**2**	**1**	0	**8**
Records that say it happened a fix announced but never applied; tickets closed as recovered on no evidence (E1, 13, 19, 20, 49)	**2**	**3**	0	0	**5**
Data and resources decision history deleted before it was archived; disk never freed (9, 10, 15)	0	**3**	0	0	**3**
People on the receiving end the same email twice, to the lead who had just replied; a weekly report that never arrives (16, 17, 57)	**2**	**1**	0	0	**3**

Fixed upstream · Confirmed, open · Deferred, waiting on a decision · Awaiting review

Cell shade grows with the count

HOW TRUST IS EARNED

Four kinds of evidence, each checkable by someone else

28

Verified by the platform

Each capability is confirmed by the platform reading this repository, not by a self-report.

53 / 66

Held under review

Findings A.M. checked line by line that held as written, or turned out worse.

5

Public fixes

Platform pull requests 3410, 3422, 3430, 3435 and 3439 fixed fourteen findings.

566

Tests that catch defects

Key rules are broken on purpose, one at a time, and a named test is shown to fail each time.

RESIDUAL RISK

What is not yet covered, and the next control for each

RESIDUAL	NEXT CONTROL
Nothing has run on live records All evidence so far is tests, examples and real platform code, not real platform data.	One export compared with Ellis's answers, together with K.D., before any capability moves up a level.
The sandbox is a seat belt, not a security boundary Node's permission model stops trusted code making mistakes; its own documentation says it does not protect against malicious code.	Baselines run only test code Ellis's engineer wrote; anything untrusted would need a container or a separate machine.
K.D. cannot see what K.D. approves Every surface is admin-only, and K.D. has no admin login.	An admin account for K.D., or a sourced written briefing until then; a governance decision.
Copied platform rules can go stale The platform changes daily; a rule Ellis copied may no longer match.	Re-check each capability's pinned citations against the latest platform commit before its first live run.

SOURCES

- Ellis AI Employee Charter: the hard limits and decision rights behind Part 1
- Colaberry Platform Reliability Audit: the 24 live risks behind Part 2
- Ellis source: every control named in the register
- Node.js permission model: what it restricts, and that it is not a security boundary against malicious code