Ellis Executive Brief

Where Ellis stands on 5 November 2026, what it has shown on the platform, and the decisions that would let it move further.

For K.D. (accountable) and A.M. (governance) **Prepared by** Quincy Nkwain Ninying, Forward Deployed Engineer

Platform read at `a29dfb245`

THE BOTTOM LINE

The evidence layer of the platform's reliability agent is built and proven. 39 of 48 capabilities are verified by the platform, and every sandbox rung of autonomy on the board is built, each with a control that shows its check can fail. Its findings have been worth acting on: 14 have already been fixed upstream.

What holds it at this rung now is decisions, not code. Three would unlock the most: K.D.'s severity table, K.D.'s thresholds for shadow mode, and A.M. naming who builds the gate and tool broker that let the checks run as an agent.

39 / 48	**16** / 16	**61** / 76	**14**	**883**
capabilities verified by the platform	sandbox rungs (R3, R4) built, each with a control that can fail	findings A.M. reviewed that held as written or worse	findings the platform has fixed upstream	tests, each key rule broken on purpose and caught

WHAT ELLIS IS

A reliability agent, built in phases, that reads and never writes

Ellis answers K.D.'s reliability questions about the Colaberry platform from its source and records, and says plainly when it cannot see. This phase is the agent's evidence and judgement layer; model-backed reasoning comes later, under the same rules.

What it does

- **Checks** the platform's jobs, timers, registry, access and recovery, one capability at a time
- **Answers** in one shape: a verdict, the rows it rests on, and what it could not see
- **Rehearses** fixes and actions on sandbox copies, beside a control that must fail
- **Reports** to K.D., who decides

What it is built to be unable to do

- **Write** to the platform, its database or its timers
- **Merge or deploy**: both are default-denied to agents; a person does both
- **Read as healthy** when it could not look: unknown is decided before clean
- **Promote itself**: K.D. signs rungs R0 to R2, A.M. signs R3 and up

WHAT IT HAS SHOWN

Findings the platform acted on, mechanisms proven in a sandbox, and answers kept honest

14 fixed

Of 82 findings, A.M. reviewed 76 and 61 held as written or worse. Among those the platform fixed:

- a scoped staff role able to raise itself to owner
- the fix agent telling reviewers a fix was applied when none was
- a lead sent the same message twice
- pause ignored by five timers

16 rungs

Each sandbox rung runs the real mechanism beside a control. What they showed:

- the preview reaper stops a stack twice when two runs overlap, despite its guard
- the registry's health score gives an agent that never ran the same 100 as a flawless one
- a backup cut cleanly between statements restores without error and answers wrongly

0 silent

Ellis corrects itself in the open:

- the platform fixed half of finding 54 on 28 October; on the day Ellis found the fix on main, the two capabilities that described it were corrected
- a sandbox rung caught a defect in Ellis's own proposal code, which is fixed
- every check names its blind spots instead of reading as clean

STILL OPEN ON THE PLATFORM

11 live risks are confirmed and open, and 3 are deferred

Three that matter most, because the platform reports the opposite of what happened. All are in the Platform Reliability Audit with their evidence.

FINDING	WHAT HAPPENS	AREA
8	Critical security alerts can never become tickets `live`	Security oversight rollup
20	A failed query, or a quiet week, closes lead-drop tickets as "cleared" `live`	Ticket auto-resolution
9	The intelligence cleanup can delete rows it never archived `live`	Data retention sweeps

DECISIONS THAT MOVE ELLIS FORWARD

Ten decisions, three of them first

Each is a person's call by design. The first three unlock the most; the rest unlock one capability each. Each row gives a recommendation to react to.

DECISION	OWNER	WHAT IT UNLOCKS	RECOMMENDATION
The severity table Which problems are Critical, High, Medium or Low (ISS-032)	`K.D.`	A severity on every problem; today each one is unassigned and says why	Start with the five first monitors' cases

DECISION	OWNER	WHAT IT UNLOCKS	RECOMMENDATION
Shadow-mode thresholds Acceptable misses and false alarms per monitor	K.D.	17 R5 rungs: running beside the platform against approved outcomes	Set them for the same five monitors first
Who builds the gate and broker And the agent core, entry points and audit log, with their host	A.M.	The checks running as an agent's tools, with reports reaching K.D.	Name one owner and a host before the next phase is planned
Rung ids in the plan	A.M.	The 16 built rungs showing as verified on the board	Republish the plan carrying rung ids
The health scoring model And what an agent with no runs should show	K.D.	A registry score that tells "never ran" from "never failed"	Adopt it, and show "unscored", not a number
Timers with no governing row Which should run, and with which enabled value	K.D.	The default-on guard and the untracked timer proposals applied	Rule on the ones that run untracked today first
Defect classes for bug-fix dispatch	K.D.	Any bug-fix rung above R1 meaning something	Begin with defects that have a failing test
The error-class vocabulary	K.D.	Failures grouped by kind across every job	Agree the platform standard's six classes
Limits per case Tool calls, time and cost	K.D.	The governance gate's budget	Confirm the limits already proposed
A backup policy Recovery objective, destination and cadence	K.D. + A.M.	Restore verification on a real backup, not a sandbox one	Set the recovery objective first; the rest follows from it

NEXT

What happens once the first three decisions land

ON THE SEVERITY TABLE

Every problem carries a severity

Ellis applies K.D.'s rules as written; no rule, no severity, and the answer says so.

ON THE THRESHOLDS

Shadow mode for five monitors

Each runs beside the platform, scored against outcomes K.D. approves, before anything reaches production.

ON AN OWNER FOR THE GATE

The checks become the agent's tools

Behind the gate and broker, under the same rules: read only, verdicts from code, a person decides.

Sources: the capability catalog and README (`docs/capabilities/README.md`); the Colaberry Platform Reliability Audit (82 findings; outcomes as of 5 Nov 2026); the platform board's rung stories; the capability docs that name each open decision; K.D.'s architecture of 28 Oct 2026. Companion documents: Ellis Problem Brief, Ellis Architecture, Ellis Workflow, Ellis Trust and Risk Review, Ellis Governance Brief.