Ellis Governance Operating Model

How the work was governed in practice: the review loop with the platform's owners, the rules it produced, the gates every change passed, and the rhythm that kept it moving.

As of 1 Nov 2026 **Prepared by** Quincy Nkwain Ninying, Forward Deployed Engineer **Accountable** K.D. **Governance** A.M.

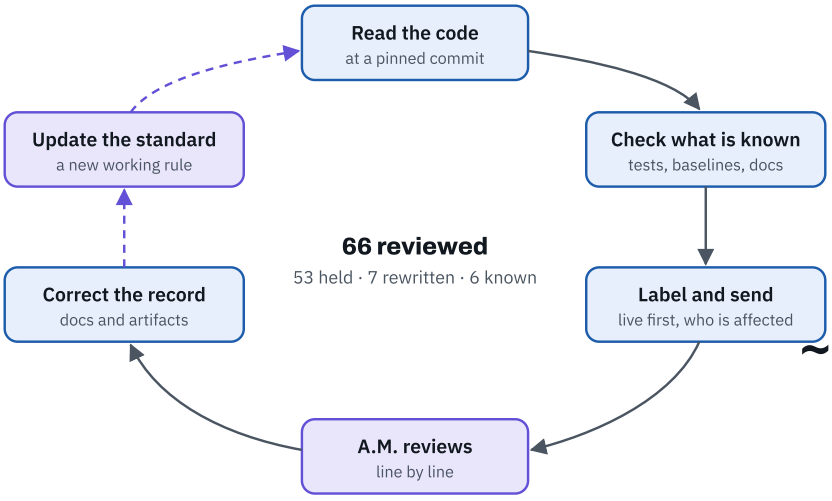
Period 19 to 30 Oct 2026

16	42	66	7	0
working rules adopted in twelve days, each from a named decision	stories verified by the platform: 6 discovery, 36 capabilities	findings reviewed line by line by A.M.	findings rewritten after review, each correction published	production actions taken by Ellis

THE REVIEW LOOP

Every finding went round the same loop, and A.M.'s corrections changed the loop

Purple marks A.M.'s steps. The dashed return is the part that matters: a correction does not only fix one finding, it becomes a rule for every later one. The emailed loop ended on 30 Oct, when K.D. and A.M. decided defects are for Ellis to find; its rules carry over to what Ellis reports.



8 working rules A.M. added through the loop: send security findings at once, write down the no-evidence rule, label every finding, check what is known first, how artifacts are published, ask the person when a finding turns on what reached them, write every verdict down per finding, and send pages as HTML or PDF

~1 day from sending a batch to A.M.'s line-by-line reply, for each batch reviewed so far

14 findings fixed in the platform through five public pull requests, 3410, 3422, 3430, 3435 and 3439

RULES AS THEY WERE ADOPTED

Sixteen working rules, each from a named decision

Oldest first. Each one still applies to every later story and artifact.

K.D. accountable manager A.M. governance Q.N. Forward Deployed Engineer

19 Oct — **K.D. Build the agent first; measure before adding fixed steps**
Deterministic steps are added only where measurements show the agent is inefficient. Discovery story 1, version 2.

21 Oct — **K.D. Some work moves to the agent, some is shared, 23 decisions stay human**
The split every capability is built inside. Discovery story 3, version 2.

22 Oct — **K.D. Argue from the platform's own figures**
The system class was approved only once it was argued from the domain's real numbers, with sources. Discovery story 1, version 3.

22 Oct — **Q.N. The public repository holds approved deliverables only**
No credentials, customer data or private notes; discovery documents go in only in the version K.D. approved.

26 Oct — **K.D. Build with what exists and demo first**
Do not wait for every document's sign-off; A.M. refines after a demo. The three proposed tasks proceed in order.

26 Oct — **A.M. Security-shaped findings go straight away, K.D. copied**
Never batched with other findings.

27 Oct — **A.M. Write down "no evidence is not health" before building more**
One page for the next person who builds a monitor: a check says healthy only on recent, positive evidence.

27 Oct — **A.M. Label every finding, live risks first, with who is affected**
Live risk, latent, or reporting gap. Overstating one finding costs trust in the next.

27 Oct — **A.M. Check what is already known before calling it new**
Search the platform's tests, baselines and documents; cite a known item instead of pitching it.

27 Oct — **A.M. Artifacts are designed pages, kept in the repository as PDF**
Built as HTML with charts and conditional formatting, published as a claude.ai copy, and only the PDF committed.

27 Oct — **A.M. When a finding turns on what reached a person, ask the person**
The code showed the Sunday report could not run; A.M. confirmed it had not arrived in 45 days.

27 Oct — **A.M. A verdict that is not written down is not a verdict**
Every review outcome is recorded per finding, with its evidence, before it goes into a commit or a report.

28 Oct — **K.D. Checks become agent tools behind a gate and a broker**
Answers carry a verdict, a severity and row ids. Where capabilities overlap, A.M. decides, and holds stand until then.

28 Oct **A.M.** **Send pages as HTML or PDF, not artifact links**
Anything sent to K.D. or A.M. is rendered from source.

30 Oct **K.D. + A.M.** **Defects are for Ellis to find, not to be sent for fixing**
Defect emails stop. Defects appear in Ellis's output and write-ups.

30 Oct **A.M.** **A release is a rung of autonomy, not a phase of work**
K.D. signs R0 to R2, A.M. signs R3 and above. Finishing a rung promotes nobody; a named person does, one step, on evidence.

APPROVAL GATES

Six gates, each with a named owner and a proof

Discovery documents

OWNER · K.D.

A version is approved before it enters the repository.

Proof: each document's version history names the approved version

Each capability

OWNER · THE PLATFORM

Built with tests, key rules broken on purpose, then confirmed by the platform reading the repository.

Proof: a verified state in the platform's progress record

Each finding

OWNER · A.M.

Until 30 Oct, reviewed line by line; any correction goes into the capability write-up and the audit.

Proof: the audit's outcome column

Each step up in autonomy

OWNER · K.D. TO R2, A.M. FROM R3

One level at a time, on evidence named in advance.

Proof: none taken yet; every capability is at its first release

Any production action

OWNER · A PERSON

Merging, deploying or switching an agent is always done by a person.

Proof: the platform denies these to agents (executionPolicy.ts:33-42)

Every push

OWNER · Q.N.

Nothing is committed or pushed without an explicit go, after a scan for anything private.

Proof: the repository's commit history

WORKING RHYTHM

How information moves, and how fast

WHAT	TO WHOM	WHEN
Security-shaped finding	A.M., K.D. copied	Same day, on its own
Live risk	A.M. and K.D.	Next message, listed first
Reporting gaps and latent defects	A.M.	Batched per capability
Capability built and verified	The platform, then the README	Each capability
Progress, blockers and decisions needed	K.D.	Weekly, Mondays 1:30 pm Central (proposed)
Artifacts	The repository (PDF) and claude.ai	Updated whenever a number they show changes

From 30 Oct, the three finding rows carry what Ellis reports. Defects are no longer hunted and emailed by hand.

SOURCES

- Discovery stories 1 to 5: the approved versions and their version histories
- Colaberry Platform Reliability Audit: review outcomes for every finding
- No evidence is not health: the rule A.M. asked for
- Ellis AI Employee Charter: decision rights, hard limits and escalation